

TRAINABLE DOWNSTREAM MODEL

FROZEN FEATURE EXTRACTOR

Utterance text
Video frames

Frozen Qwen Thinker
last-token features

raw multimodal input

L2 normalization
3584→128 projection

Inherited DGF
normal path

JS-DG robust path
aux. disagreement

EXTENSION UNDER AUDIT

g

Five-turn context
speaker + distance

SPM read
global fallback

Train-only speaker
prototype library

64-32-7 head
main + aux. CE

seven-class prediction

